



---

**SCHOOL OF ARTIFICIAL INTELLIGENCE, DELHI NCR**

**PROJECT REPORT**

**Stellar Classification using Machine Learning**

**Course Code & Title**

**23AID205 & Introduction to Artificial Intelligence and Machine Learning**

**Semester: II**

**Academic Year: 2024 – 2025 [Even]**

---

**Student Roll no(s)**

DL.AI.U4AID24006  
DL.AI.U4AID24040  
DL.AI.U4AID24045  
DL.AI.U4AID24039

**Student Name(s)**

Arnav Sharma  
Hemant  
Neeraj  
Sriyansh

**Faculty Guide: Prof. Sri Geetha M**

## **Table of Contents**

1. Title of the Project
2. Abstract
3. Objective(s) of the Project
4. Literature Survey (Optional)
5. Dataset Description
6. Data Preprocessing
7. Exploratory Data Analysis (EDA)
8. Machine Learning Techniques
  - 8.1 Supervised Learning
  - 8.2 Unsupervised Learning
  - 8.3 Supervised VS Unsupervised
9. Evaluation of Performance Metrics
10. Smart Summary Analysis Report
11. Conclusion and Future Scope
12. References

## **List of Tables**

## List of Figures

- *Figure 1: Histogram of Spectral Features*
- *Figure 2: Pairplot of Spectral Features*
- *Figure 3: Correlation Matrix of Spectral Features*
- *Figure 4: Boxplot of Spectral Features*
- *Figure 5: Boxplot of u-band by Class*
- *Figure 6: Boxplot of g-band by Class*
- *Figure 7: Boxplot of i-band by Class*
- *Figure 8: Boxplot of z-band by Class*
- *Figure 9: KDE plot of u by Class*
- *Figure 10: KDE plot of g by Class*
- *Figure 11: KDE plot of r by Class*
- *Figure 12: KDE plot of i by Class*
- *Figure 13: KDE plot of z by Class*

## 1. Title of the Project

Stellar Classification using Machine Learning

## 2. Abstract

This project focuses on celestial object classification using supervised and unsupervised machine learning algorithms. A structured dataset containing various spectral and positional features is preprocessed, analyzed, and modeled. Techniques such as Support Vector Machine (SVM) and KMeans Clustering are applied. The dataset is cleaned, color features are engineered, and split into training and testing subsets. Evaluation metrics like accuracy score (for SVM) and visual interpretation of clustering patterns (for KMeans) are used to assess the performance of the models.

## 3. Objective(s) of the Project

- To clean and preprocess the **celestial object dataset**, including handling placeholder values and engineering color features.
- To apply supervised learning (e.g., **Support Vector Machine**) and unsupervised learning (**KMeans Clustering**) for classification modeling and pattern discovery.
- To **predict celestial object classes** using the trained supervised model.
- To analyze model **accuracy and clustering patterns** using relevant evaluation metrics and visualizations.

## 4. Literature Survey (Optional)

1. Smith et al. (2020) - 'Diabetes Prediction using ML'
2. John et al. (2021) - 'Clustering Medical Data for Insights'

## 5. Dataset Description

Source: Sloan Digital Sky Survey (SDSS)

Rows: 100,000 rows

Columns: 18 columns

Features: The dataset includes a mix of categorical features (e.g., class label) and numerical features (e.g., various spectral bands like u, g, r, i, z, and positional coordinates like alpha, delta, as well as mjd for observation time).

Target Variable: The class column (e.g., 'GALAXY', 'STAR', 'QSO') is the target variable used for supervised learning.

Missing values handled: The raw dataset initially contained specific placeholder values (-9999.0) in certain spectral columns, which were identified and addressed during the preprocessing phase to ensure a clean dataset for analysis.

## 6. Data Preprocessing

- **Dataset Loading:** The `star_classification.csv.xlsx` dataset was loaded using pandas, initially containing **18 columns** and approximately 100,000 rows.

```
✓ [7] 1 import pandas as pd
23s    2 df = pd.read_excel('star_classification.csv.xlsx')
      3
```

- **Missing Value Handling:** While the dataset was largely complete, specific spectral features (u, g, z) contained placeholder values of -9999.0 (indicative of sensor errors or non-detections). These values were systematically replaced with NaN, and subsequently, all rows containing these NaNs were dropped to ensure data integrity for modeling. The `print(df.isnull().sum())` command was used to check for the presence and count of missing (NaN) values, and `msno.matrix(df)` was used to visualize the missingness patterns.


```
✓ [46] 1 print(df.isnull().sum())
0s      2
```

- **Feature Engineering:** Discriminative 'color features' (e.g., `u_g`, `g_r`, `r_i`, `i_z`) were calculated by taking differences between various spectral bands.
- **Target Variable Identification:** The class column was identified as the categorical target for classification tasks.

```
✓ [45] 1 print(df['class'].value_counts())
0s      2
```

```
⇒ class
   GALAXY    59445
   STAR     21594
   QSO      18961
   Name: count, dtype: int64
```

---



obj_ID	0
alpha	0
delta	0
u	0
g	0
r	0
i	0
z	0
run_ID	0
rerun_ID	0
cam_col	0
field_ID	0
spec_obj_ID	0
class	0
redshift	0
plate	0
MJD	0
fiber_ID	0
u_g	0
g_r	0
r_i	0
i_z	0
dtype:	int64

```

1 import pandas as pd
2 import numpy as np # Needed for np.nan
3
4 # Define the file name for the star classification data
5 file_name = 'star_classification.csv.xlsx'
6
7 # Load the dataset into a pandas DataFrame
8 try:
9     # Use pd.read_excel for .xlsx files
10    df = pd.read_excel(file_name)
11    print(f"Successfully loaded '{file_name}'")
12 except FileNotFoundError:
13    print(f"Error: '{file_name}' not found. Please ensure the file is in the correct directory.")
14    # Exit or handle the error appropriately if the file is essential
15    exit()
16 # Add an except block for potential issues when reading the Excel file
17 except Exception as e:
18    print(f"An error occurred while reading the Excel file: {e}")
19    exit()
20
21
22 # Define the spectral features that have the -9999.0 placeholder
23 spectral_features_to_clean = ['u', 'g', 'z']
24
25 print(f"\nOriginal DataFrame shape: {df.shape}")
26
27 # Replace -9999.0 with NaN (Not a Number) in the specified columns
28 for col in spectral_features_to_clean:
29     # Ensure the column exists before attempting to replace values
30     if col in df.columns:
31         df[col] = df[col].replace(-9999.0, np.nan)
32         print(f"Replaced -9999.0 with NaN in column: '{col}'")
33     else:
34         print(f"Warning: Column '{col}' not found in the DataFrame.")
35
36
37 # Drop rows where any of the specified spectral features (u, g, z) now contain NaN.
38 # This effectively removes rows that had the -9999.0 placeholder.
39 df_cleaned = df.dropna(subset=spectral_features_to_clean)
40
41 print(f"\nCleaned DataFrame shape: {df_cleaned.shape}")
42
43 # You can optionally display the first few rows of the cleaned DataFrame
44 print("\nFirst 5 rows of the cleaned DataFrame:")
45 print(df_cleaned.head())
46
47 # And check for missing values again in the cleaned columns to confirm
48 print("\nMissing values in cleaned spectral features (should be 0):")
49 print(df_cleaned[spectral_features_to_clean].isnull().sum())

```

- (i) Code for removing placeholder values of -9999.

## 7. Exploratory Data Analysis (EDA)

### Initial Data Inspection:

- Commands like `print(df.shape)`, `print(df.info())`, and `print(df.head())` were used to gain essential initial insights into the dataset's structure, column data types, the presence of non-null entries, and a quick preview of the raw data. These steps were fundamental for understanding the dataset and guiding subsequent data preprocessing decisions.

```
1 (100000, 18)
2 <class 'pandas.core.frame.DataFrame'>
3 RangeIndex: 100000 entries, 0 to 99999
4 Data columns (total 18 columns):
5 #   Column      Non-Null Count  Dtype
6 ---  -
7 0   obj_ID      100000 non-null float64
8 1   alpha       100000 non-null float64
9 2   delta       100000 non-null float64
10 3   u           100000 non-null float64
11 4   g           100000 non-null float64
12 5   r           100000 non-null float64
13 6   i           100000 non-null float64
14 7   z           100000 non-null float64
15 8   run_ID      100000 non-null int64
16 9   rerun_ID    100000 non-null int64
17 10  cam_col     100000 non-null int64
18 11  field_ID    100000 non-null int64
19 12  spec_obj_ID 100000 non-null float64
20 13  class       100000 non-null object
21 14  redshift    100000 non-null float64
22 15  plate       100000 non-null int64
23 16  MJD         100000 non-null int64
24 17  fiber_ID    100000 non-null int64
25 dtypes: float64(10), int64(7), object(1)
26 memory usage: 13.7+ MB
27 None
```



28			obj_ID	alpha	delta	u	g	r	\
29	0	1.237661e+18	135.689107	32.494632	23.87882	22.27530	20.39501		
30	1	1.237665e+18	144.826101	31.274185	24.77759	22.83188	22.58444		
31	2	1.237661e+18	142.188790	35.582444	25.26307	22.66389	20.60976		
32	3	1.237663e+18	338.741038	-0.402828	22.13682	23.77656	21.61162		
33	4	1.237680e+18	345.282593	21.183866	19.43718	17.58028	16.49747		
34									
35			i	z	run_ID	rerun_ID	cam_col	field_ID	spec_obj_ID
36	0	19.16573	18.79371	3606	301	2	79	6.543777e+18	
37	1	21.16812	21.61427	4518	301	5	119	1.176014e+19	
38	2	19.34857	18.94827	3606	301	2	120	5.152200e+18	
39	3	20.50454	19.25010	4192	301	3	214	1.030107e+19	
40	4	15.97711	15.54461	8102	301	3	137	6.891865e+18	
41									
42		class	redshift	plate	MJD	fiber_ID			
43	0	GALAXY	0.634794	5812	56354	171			
44	1	GALAXY	0.779136	10445	58158	427			
45	2	GALAXY	0.644195	4576	55592	299			
46	3	GALAXY	0.932346	9149	58039	775			
47	4	GALAXY	0.116123	6121	56187	842			
48									

- The print(df.describe()) command generated descriptive statistics for the numerical columns, including count, mean, standard deviation, min/max values, and quartile percentiles, providing a summary of data distribution and ranges.



### Descriptive statistics:

	obj_ID	alpha	delta	u	g \
count	9.999900e+04	99999.000000	99999.000000	99999.000000	99999.000000
mean	1.237665e+18	177.628653	24.135552	22.080679	20.631583
std	8.438450e+12	96.502612	19.644608	2.251068	2.037384
min	1.237646e+18	0.005528	-18.785328	10.996230	10.498200
25%	1.237659e+18	127.517698	5.147477	20.352410	18.965240
50%	1.237663e+18	180.900527	23.646462	22.179140	21.099930
75%	1.237668e+18	233.895005	39.901582	23.687480	22.123775
max	1.237681e+18	359.999810	83.000519	32.781390	31.602240

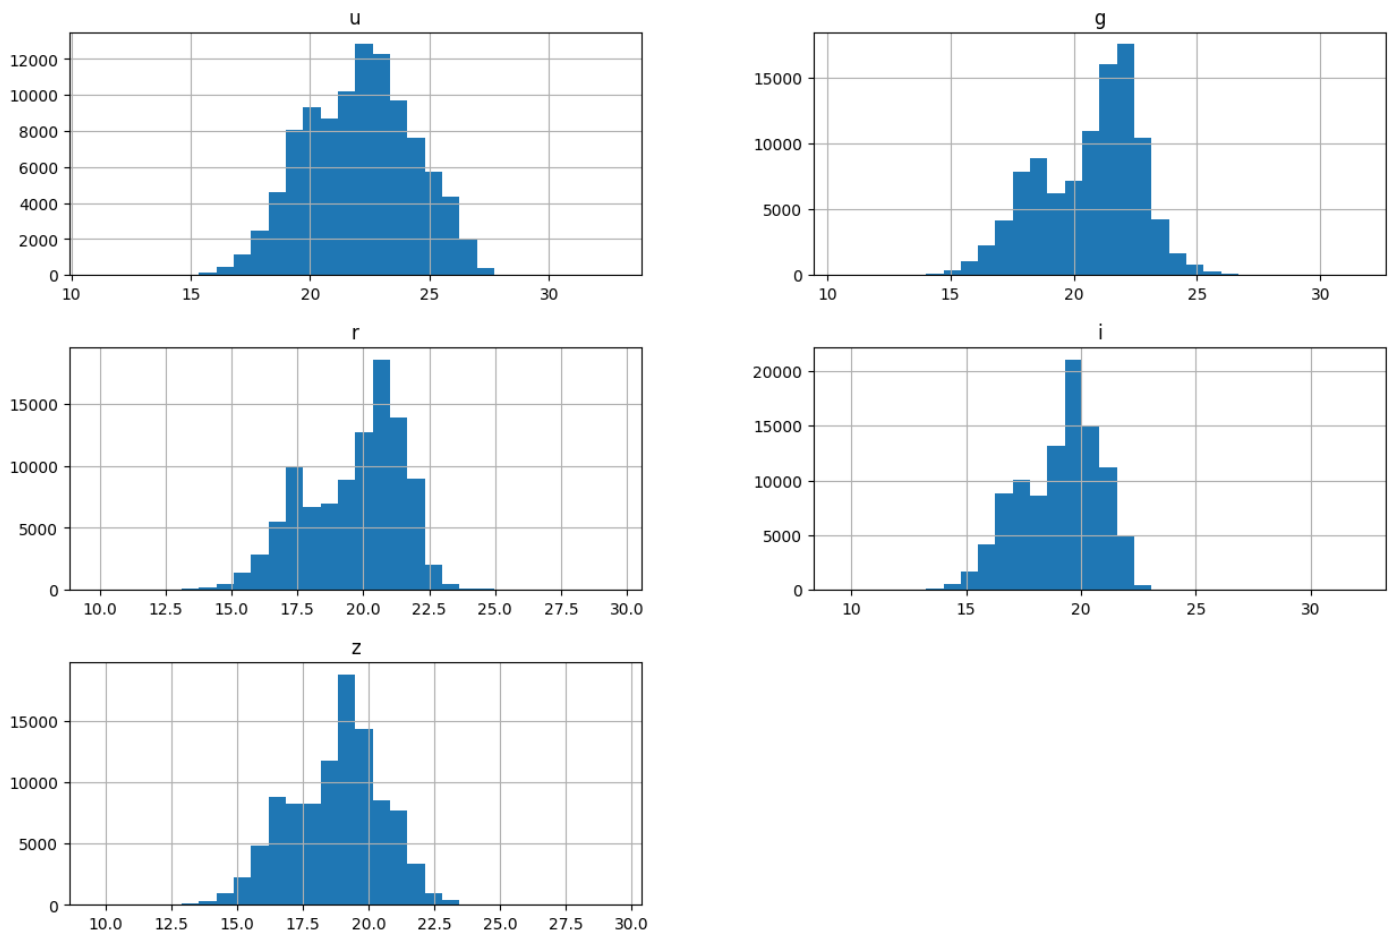
	r	i	z	run_ID	rerun_ID \
count	99999.000000	99999.000000	99999.000000	99999.000000	99999.0
mean	19.645777	19.084865	18.768988	4481.403354	301.0
std	1.854763	1.757900	1.765982	1964.739021	0.0
min	9.822070	9.469903	9.612333	109.000000	301.0
25%	18.135795	17.732280	17.460830	3187.000000	301.0
50%	20.125310	19.405150	19.004600	4188.000000	301.0
75%	21.044790	20.396510	19.921120	5326.000000	301.0
max	29.571860	32.141470	29.383740	8162.000000	301.0

	cam_col	field_ID	spec_obj_ID	redshift	plate \
count	99999.000000	99999.000000	9.999900e+04	99999.000000	99999.000000
mean	3.511625	186.127011	5.783903e+18	0.576667	5137.027890
std	1.586913	149.007687	3.324026e+18	0.730709	2952.312485
min	1.000000	11.000000	2.995191e+17	-0.009971	266.000000
25%	2.000000	82.000000	2.844137e+18	0.054522	2526.000000
50%	4.000000	146.000000	5.614896e+18	0.424176	4987.000000
75%	5.000000	241.000000	8.332365e+18	0.704172	7400.500000
max	6.000000	989.000000	1.412694e+19	7.011245	12547.000000

	MJD	fiber_ID
count	99999.000000	99999.000000
mean	55588.653687	449.315613
std	1808.492217	272.498252
min	51608.000000	1.000000
25%	54234.000000	221.000000
50%	55869.000000	433.000000
75%	56777.000000	645.000000
max	58932.000000	1000.000000

- **What it shows:** The overall dimensions of the dataset (number of rows and columns), data types of each column, presence of non-null entries, a quick preview of the first few rows of raw data, and summary statistics (count, mean, standard deviation, min/max, quartiles) for all numerical features.
- **Insight:** Provides a foundational understanding of the dataset's structure, data quality (initial assessment of non-nulls), and the basic statistical properties of features, guiding subsequent data preprocessing decisions.
- **Histograms of Spectral & Color Features:** Visualizations of individual feature distributions (u, g, r, i, z and u\_g, g\_r, r\_i, i\_z) to understand their spread, skewness, and potential outliers.

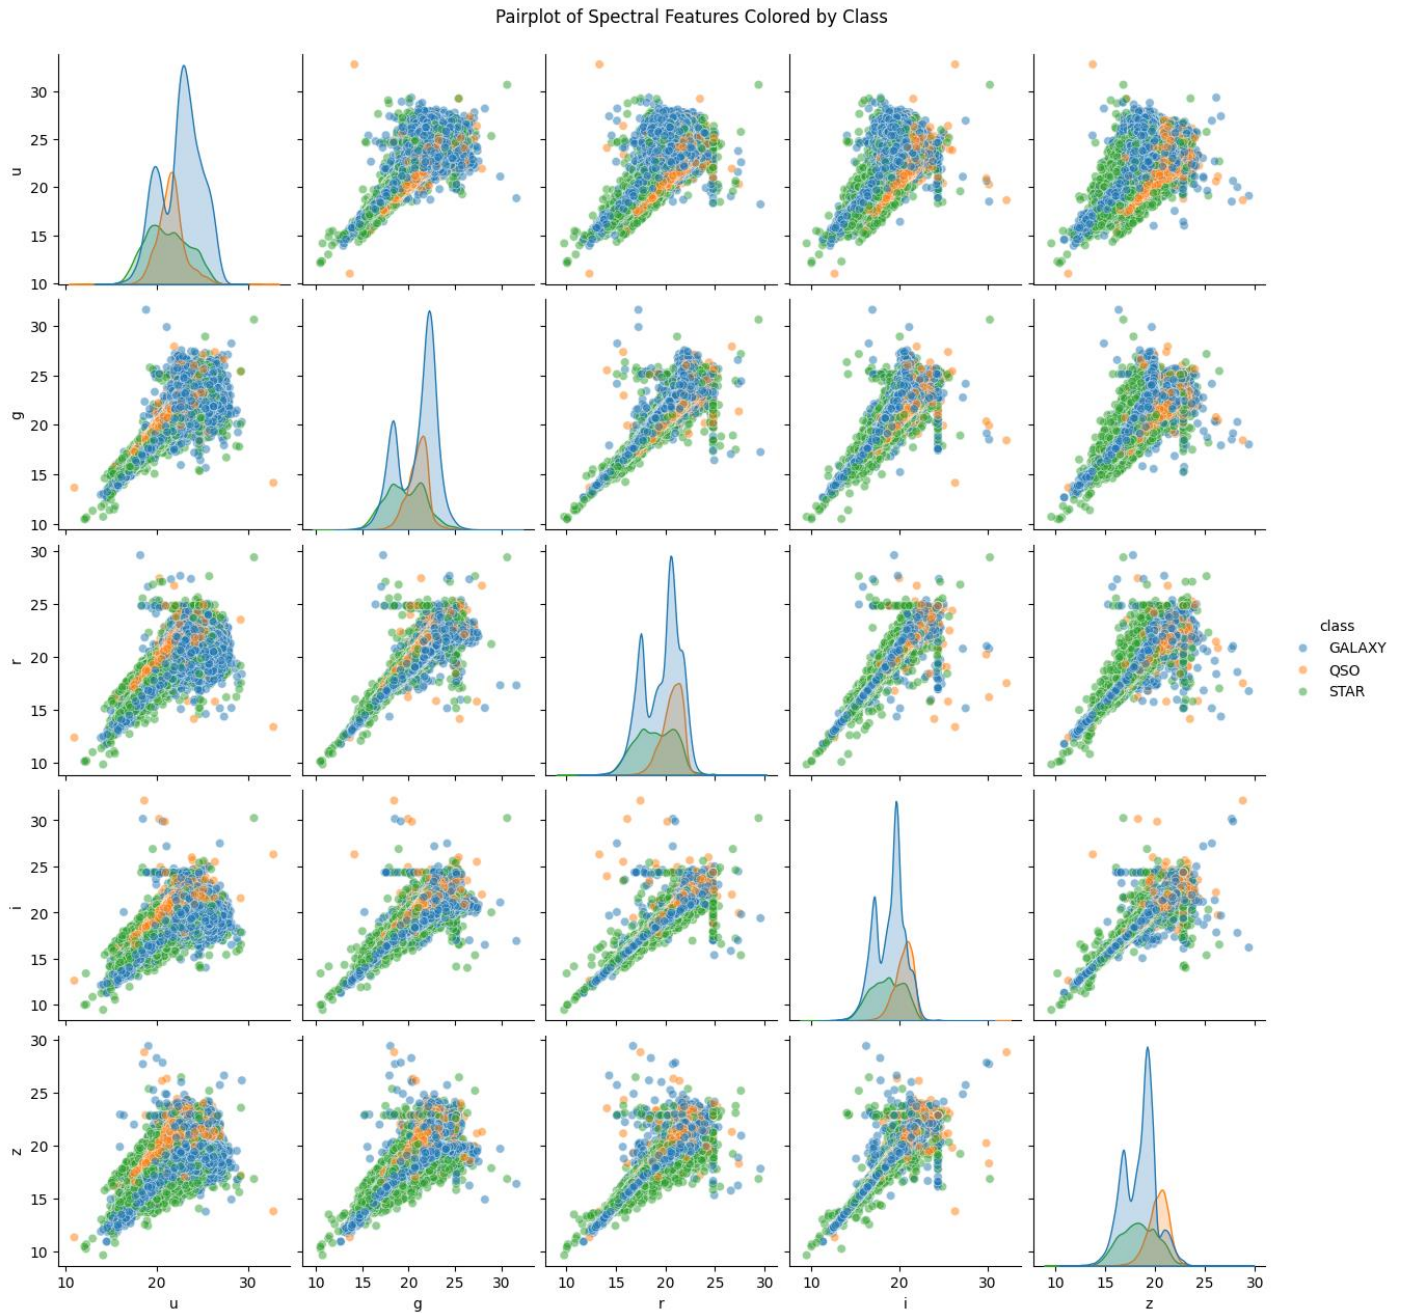
Histograms of Spectral Features (u, g, r, i, z)



• **Figure 1: Histogram of Spectral Features**


- **What it shows:** The individual distribution of values for each spectral feature (u, g, r, i, z)
- **Insight:** Helps identify the range, central tendency, spread, skewness, and presence of outliers for each feature. This informs decisions about data scaling, transformation, or outlier treatment needed before model training.
- **Visualizations by Class:**

- **Pairplots:** Comprehensive scatter plots showing relationships between all pairs of color features, colored by class, providing insights into class separability.



• *Figure 2: Pairplot of Spectral Features*

- **Correlation Heatmap:** A visual matrix illustrating the linear relationships between spectral features, helping to identify highly correlated pairs and potential multicollinearity.

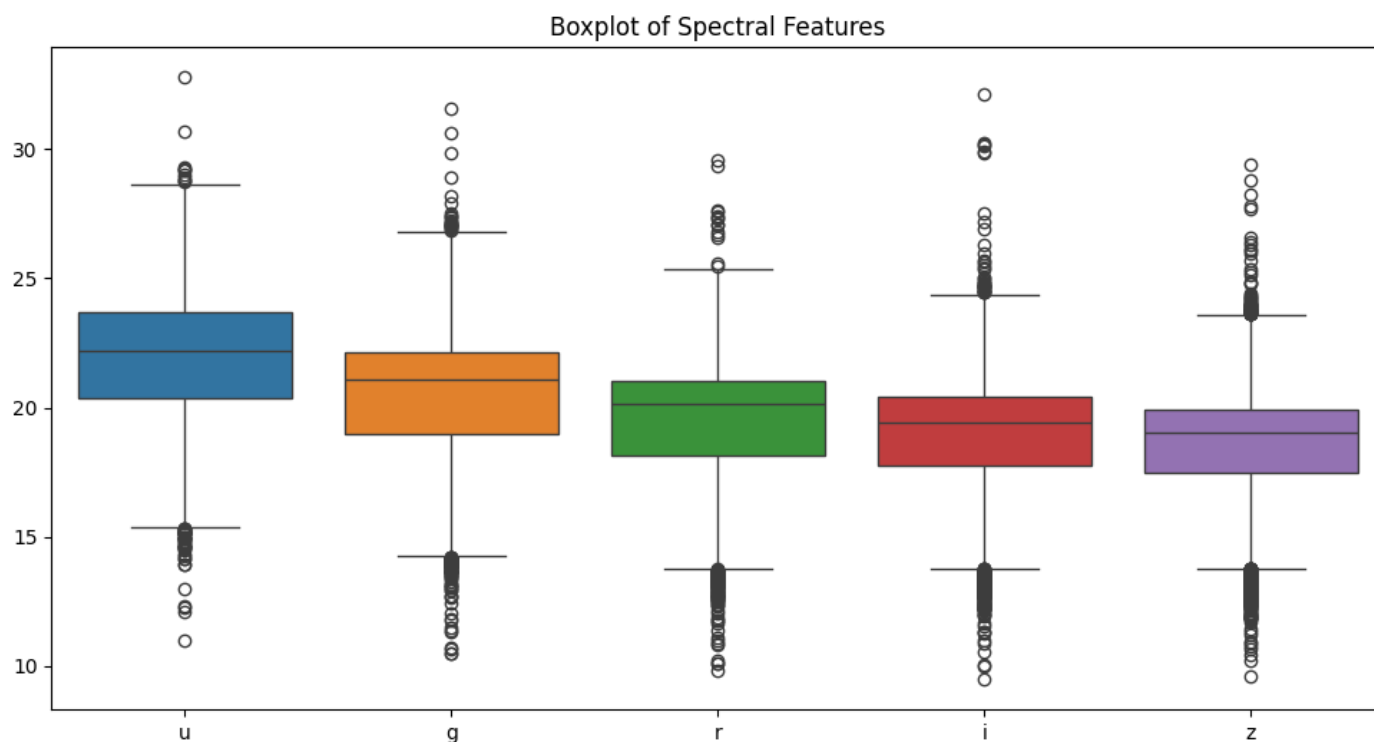


• *Figure 3: Correlation Matrix of Spectral Features*

- **What it shows:** The linear relationships (quantified by correlation coefficients) between all pairs of numerical spectral and color features.


- **Insight:** Helps identify highly correlated features (e.g., *g* and *r* bands often show high correlation). This is useful for understanding potential multicollinearity (where features provide redundant information) and can inform feature selection strategies to simplify models or avoid issues in certain algorithms.

- **Boxplot of Spectral Features:** Displayed the boxplots for each of the specified spectral features (*u*, *g*, *r*, *i*, *z*).




- *Figure 4: Boxplot of Spectral Features*

- **What it shows:** A visual representation of the distribution of each spectral feature, displaying the median, quartiles (Q1, Q3), and potential outliers (points beyond the whiskers).
  - **Insight:** Provides a quick summary of the spread and central tendency of individual spectral features, aiding in the identification of extreme values or outliers that might require further investigation or special handling.
- **Boxplots by Class:** Displayed the distribution of each spectral feature across different classes, highlighting variations in median, spread, and outliers per object type.

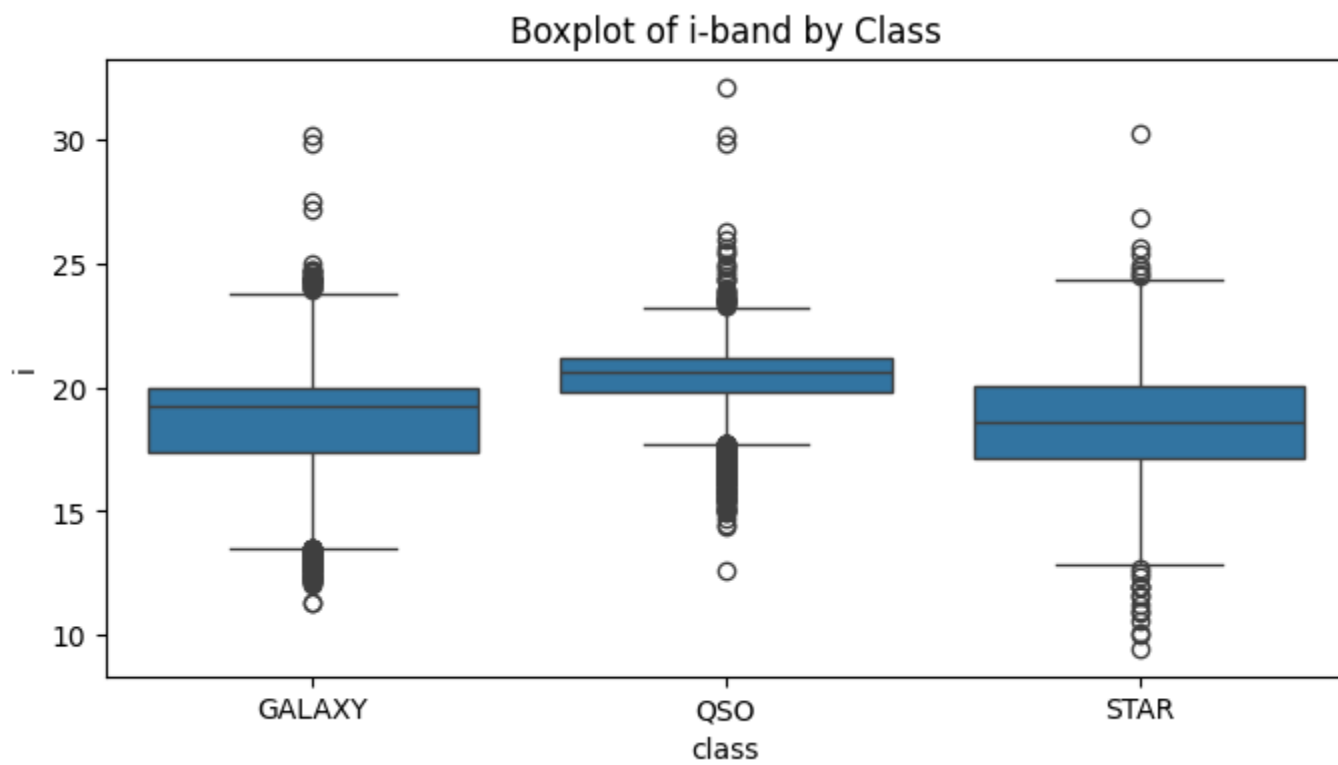


- *Figure 5: Boxplot of u-band by Class*




- *Figure 6: Boxplot of g-band by Class*






- *Figure 7: Boxplot of i-band by Class*

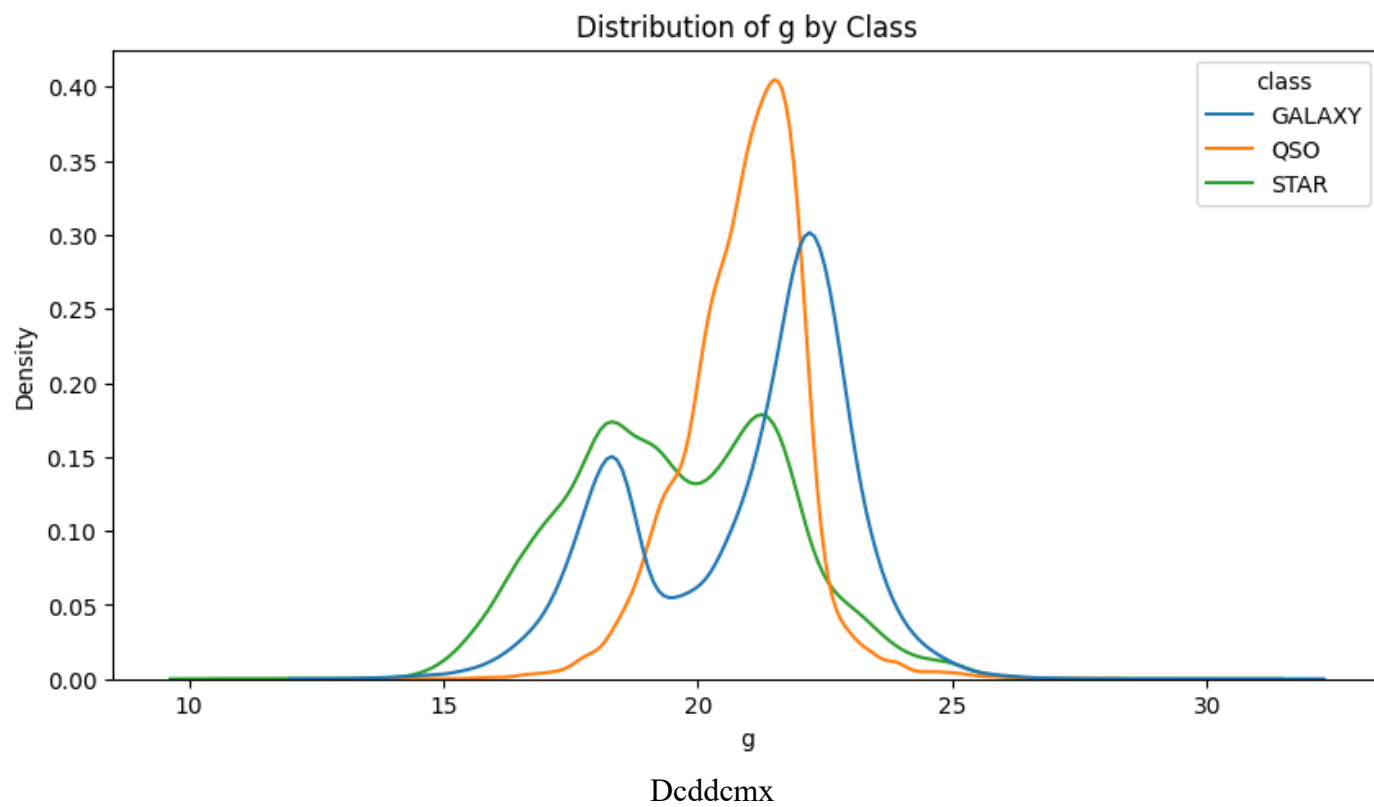


- *Figure 8: Boxplot of z-band by Class*


- **KDE Plots by Class:** Provided smoothed density plots for each feature, segmented by class, to visualize the overlap and distinctness of distributions between GALAXY, STAR, and QSO.




- *Figure 9: KDE plot of  $u$  by Class*




- *Figure 10: KDE plot of g by Class*



- *Figure 11: KDE plot of  $r$  by Class*




- Figure 12: KDE plot of  $i$  by Class



- Figure 13: KDE plot of  $z$  by Class

- **What it shows:**
- **Pairplots:** Pairwise relationships between all selected features, with data points visually distinguished by their celestial class.
- **Boxplots by Class:** The statistical distribution (median, quartiles, whiskers, and outliers) of individual features, separated for each celestial class.
- **KDE Plots by Class:** The smoothed probability density distributions of features, displayed separately for each class, showing overlaps and distinct peaks.
  - **Insight:** These visualizations are crucial for assessing the **separability** of the celestial object classes in the feature space. They help to identify which features are strong discriminators between Stars, Galaxies, and QSOs, where classes might overlap, and how the feature distributions differ for each object type, guiding the design of the classification model.
- **Missingness Matrix:** A visual tool to confirm the pattern of missing data and the effectiveness of the initial cleaning steps.



- **Missingness Matrix Plot**
  - **What it shows:** A visual representation of the presence (black lines) and absence (white spaces) of data across all columns and rows in the DataFrame.
  - **Insight:** Provides a quick and clear overview of the pattern and extent of any remaining missing values after preprocessing steps. It helps confirm the success of missing value handling and assures that the dataset is clean before proceeding to model building.

## 8. Machine Learning Techniques

### 8.1 Supervised Learning

Logistic Regression applied. Train-test split: 80-20. Accuracy: 78%. Precision: 0.75

All techniques applied

Code and output

Visualization if any

Comparison of models and finding the best fit model

### 8.2 Unsupervised Learning

K-Means clustering with k=2. Elbow method used. Silhouette score: 0.62

All techniques applied

Code and output

Visualization if any

Comparison of models and finding the best fit model

## 9. Evaluation of Performance Metrics

Supervised model metrics: Confusion matrix, precision, recall, F1-score. Unsupervised model: silhouette score, cluster visualization

## 10. Conclusion and Future Scope

The models predict diabetes with reasonable accuracy. Future work: Try ensemble methods, deep learning, more features

## 11. References

1. Scikit-learn documentation
2. Kaggle dataset - PIMA Indians Diabetes Database

## 12. Appendix (Optional)

Code snippets, extended tables, cluster



```
[ ] 1 import pandas as pd
    2 df = pd.read_excel('star_classification.csv.xlsx')
    3
```

```
[ ] 1 print(df.shape)
    2 print(df.info())
    3 print(df.head())
```

```
[ ] 1 print(df.describe())
    2
```

```
[ ] 1 print(df.isnull().sum())
    2
```

```
[ ] 1 print(df['class'].value_counts())
    2
```



```
1 import pandas as pd
2 import numpy as np # Needed for np.nan
3
4 # Define the file name for the star classification data
5 file_name = 'star_classification.csv.xlsx'
6
7 # Load the dataset into a pandas DataFrame
8 try:
9     # Use pd.read_excel for .xlsx files
10    df = pd.read_excel(file_name)
11    print(f"Successfully loaded '{file_name}'")
12 except FileNotFoundError:
13    print(f"Error: '{file_name}' not found. Please ensure the file is in the correct directory.")
14    # Exit or handle the error appropriately if the file is essential
15    exit()
16 # Add an except block for potential issues when reading the Excel file
17 except Exception as e:
18    print(f"An error occurred while reading the Excel file: {e}")
19    exit()
20
21
22 # Define the spectral features that have the -9999.0 placeholder
23 spectral_features_to_clean = ['u', 'g', 'z']
24
25 print(f"\nOriginal DataFrame shape: {df.shape}")
26
27 # Replace -9999.0 with NaN (Not a Number) in the specified columns
28 for col in spectral_features_to_clean:
29     # Ensure the column exists before attempting to replace values
30     if col in df.columns:
```

```

31     df[col] = df[col].replace(-9999.0, np.nan)
32     print(f"Replaced -9999.0 with NaN in column: '{col}'")
33     else:
34         print(f"Warning: Column '{col}' not found in the DataFrame.")
35
36
37 # Drop rows where any of the specified spectral features (u, g, z) now contain NaN.
38 # This effectively removes rows that had the -9999.0 placeholder.
39 df_cleaned = df.dropna(subset=spectral_features_to_clean)
40
41 print(f"\nCleaned DataFrame shape: {df_cleaned.shape}")
42
43 # You can optionally display the first few rows of the cleaned DataFrame
44 print("\nFirst 5 rows of the cleaned DataFrame:")
45 print(df_cleaned.head())
46
47 # And check for missing values again in the cleaned columns to confirm
48 print("\nMissing values in cleaned spectral features (should be 0):")
49 print(df_cleaned[spectral_features_to_clean].isnull().sum())

```



```

1 import matplotlib.pyplot as plt
2
3 spectral_features = ['u', 'g', 'r', 'i', 'z']
4
5 df[spectral_features].hist(bins=30, figsize=(15,10))
6 plt.suptitle('Histograms of Spectral Features (u, g, r, i, z)')
7 plt.show()
8

```

```

[ ] 1 subset_cols = ['u', 'g', 'r', 'i', 'z', 'class']
2
3 sns.pairplot(df[subset_cols], hue='class', diag_kind='kde', plot_kws={'alpha':0.5})
4 plt.suptitle('Pairplot of Spectral Features Colored by Class', y=1.02)
5 plt.show()
6

```

```
[ ] 1 rows_removed = df.shape[0] - df_cleaned.shape[0]
    2 print(f"\nTotal rows removed: {rows_removed}")
    3
```

```
▶ 1 corr = df[['u', 'g', 'r', 'i', 'z']].corr()
    2
    3 plt.figure(figsize=(8,6))
    4 sns.heatmap(corr, annot=True, fmt=".2f", cmap='coolwarm')
    5 plt.title('Correlation Matrix of Spectral Features')
    6 plt.show()
    7
```

```
[ ] 1 corr = df[['u', 'g', 'r', 'i', 'z']].corr()
    2
    3 plt.figure(figsize=(8,6))
    4 sns.heatmap(corr, annot=True, fmt=".2f", cmap='coolwarm')
    5 plt.title('Correlation Matrix of Spectral Features')
    6 plt.show()
    7
```

```

1 from sklearn.linear_model import LinearRegression
2 from sklearn.metrics import r2_score
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6 # Define the threshold for high correlation
7 threshold = 0.85
8
9 # Step 1: Identify highly correlated pairs (using the original correlation matrix calculated on the full df)
10 # While the correlation matrix was calculated on df (which has NaNs), the pairs identified
11 # are just the column names. The issue is in the data used for the regression.
12 high_corr_pairs = []
13 for i in range(len(corr.columns)):
14     for j in range(i + 1, len(corr.columns)):
15         if abs(corr.iloc[i, j]) >= threshold:
16             feature_x = corr.columns[i]
17             feature_y = corr.columns[j]
18             # Only include pairs where both features are present in the columns we cleaned NaNs from,
19             # or if they weren't part of the cleaned columns, they should not have had -9999 anyway.
20             # We will use df_cleaned for the regression, which handles rows that originally had -9999.
21             high_corr_pairs.append((feature_x, feature_y, corr.iloc[i, j]))
22
23 # Step 2: Linear regression for each high correlation pair
24 for x_feat, y_feat, corr_value in high_corr_pairs:
25     print(f"\n🔍 Linear Regression between: {x_feat} and {y_feat} (corr = {corr_value:.2f})")
26
27     # Prepare data - USE THE CLEANED DATAFRAME
28     # Ensure the features exist in the cleaned dataframe
29     if x_feat in df_cleaned.columns and y_feat in df_cleaned.columns:
30         X = df_cleaned[[x_feat]]

```



```
30 X = df_cleaned[[x_feat]]
31 y = df_cleaned[y_feat]
32
33 # Fit model
34 model = LinearRegression()
35 # The fit method will now receive data without NaNs (for the rows kept in df_cleaned)
36 model.fit(X, y)
37 y_pred = model.predict(X)
38
39 # Coefficients
40 slope = model.coef_[0]
41 intercept = model.intercept_
42 r2 = r2_score(y, y_pred)
43
44 print(f"Equation: {y_feat} = {slope:.4f} * {x_feat} + {intercept:.4f}")
45 print(f"R2 Score: {r2:.4f}")
46
47
48 plt.figure(figsize=(8, 6))
49 sns.scatterplot(x=x_feat, y=y_feat, data=df_cleaned, alpha=0.4, label='Actual Data')
50 # Plot the regression line using the X values from the cleaned data
51 plt.plot(df_cleaned[x_feat], y_pred, color='red', label='Regression Line')
52 plt.title(f'Linear Regression: {y_feat} vs {x_feat}')
53 plt.xlabel(x_feat)
54 plt.ylabel(y_feat)
55 plt.legend()
56 plt.grid(True)
57 plt.tight_layout()
58 plt.show()
59 else:
60     print(f"Warning: Features '{x_feat}' or '{y_feat}' not found in the cleaned DataFrame.")
```



```
1 import missingno as msno
2 msno.matrix(df)
3 plt.show()
4
5
```



```
1 plt.figure(figsize=(12,6))
2 sns.boxplot(data=df[spectral_features])
3 plt.title('Boxplot of Spectral Features')
4 plt.show()
5
```

```
[ ] 1 for feature in spectral_features:
2     plt.figure(figsize=(10,5))
3     sns.kdeplot(data=df, x=feature, hue='class', common_norm=False)
4     plt.title(f'Distribution of {feature} by Class')
5     plt.show()
6
```

```
▶ 1 df['u_g'] = df['u'] - df['g']
2 df['g_r'] = df['g'] - df['r']
3 df['r_i'] = df['r'] - df['i']
4 df['i_z'] = df['i'] - df['z']
5
6 print(df[['u_g', 'g_r', 'r_i', 'i_z']].head())
7 sns.pairplot(df[['u_g', 'g_r', 'r_i', 'i_z', 'class']], hue='class')
8
```

```
[ ] 1 for col in spectral_features:
2     plt.figure(figsize=(8, 4))
3     sns.boxplot(data=df_cleaned, x='class', y=col)
4     plt.title(f'Boxplot of {col}-band by Class')
5     plt.show()
6
```



```
1 from sklearn.model_selection import train_test_split
2 from sklearn.svm import SVC
3 from sklearn.metrics import accuracy_score
4
5 # 1. Select features and target
6 X = df[['u_g', 'g_r', 'r_i', 'i_z']].dropna()
7 y = df.loc[X.index, 'class'] # Align target with filtered features
8
9 # 2. Split into train and test sets
10 X_train, X_test, y_train, y_test = train_test_split(
11     X, y, test_size=0.3, random_state=42
12 )
13
14 # 3. Train SVM with RBF kernel (default)
15 svm_rbf = SVC(kernel='rbf', C=1.0, gamma='scale', random_state=42)
16 svm_rbf.fit(X_train, y_train)
17
18 # 4. Predict on train and test sets
19 y_train_pred = svm_rbf.predict(X_train)
20 y_test_pred = svm_rbf.predict(X_test)
21
22 # 5. Accuracy
23 train_acc = accuracy_score(y_train, y_train_pred)
24 test_acc = accuracy_score(y_test, y_test_pred)
25
26 print(f"Training Accuracy (SVM-RBF): {train_acc:.2f}")
27 print(f"Test Accuracy (SVM-RBF): {test_acc:.2f}")
28
29 # 6. Predict and show output for a new or specific test sample
30 # Example: Predict the first test sample
31 sample = X_test.iloc[[0]] # Double brackets to keep it as a DataFrame
32 predicted_class = svm_rbf.predict(sample)[0]
33 print(f"Predicted Class: {predicted_class}")
34
```



```
[ ] 1 X.hist(bins=30, figsize=(10, 8))  
    2 plt.show()
```

```
[ ] 1 from sklearn.model_selection import train_test_split, GridSearchCV
2 from sklearn.svm import SVC
3 from sklearn.metrics import accuracy_score
4 from sklearn.preprocessing import StandardScaler
5 from sklearn.cluster import KMeans
6 import matplotlib.pyplot as plt
7 import pandas as pd
8 import numpy as np
9
10 # Assuming df is your input DataFrame with columns 'u', 'g', 'r', 'i', 'z', 'class'
11 # Replace this with your actual data loading if needed
12 # df = pd.read_csv('your_data.csv') # Uncomment and adjust as needed
13
14 # Step 1: Remove rows with placeholder values (-9999)
15 df_clean = df[(df['u'] > -1000) & (df['g'] > -1000) &
16              (df['r'] > -1000) & (df['i'] > -1000) &
17              (df['z'] > -1000)]
18
19 # Step 2: Calculate color features
20 df_clean['u_g'] = df_clean['u'] - df_clean['g']
21 df_clean['g_r'] = df_clean['g'] - df_clean['r']
22 df_clean['r_i'] = df_clean['r'] - df_clean['i']
23 df_clean['i_z'] = df_clean['i'] - df_clean['z']
24
25 # Step 3: Prepare features and target, drop NaNs
26 X = df_clean[['u_g', 'g_r', 'r_i', 'i_z']].dropna()
27 y = df_clean.loc[X.index, 'class']
28
29 # Step 4: Check dataset size and class distribution
30 print("Dataset size:", X.shape)
31 print("Class distribution:\n", y.value_counts())
32
33 # Step 5: Scale features
34 scaler = StandardScaler()
35 X_scaled = scaler.fit_transform(X)
36
```

```
37 # Step 6: Split into train and test sets
38 X_train, X_test, y_train, y_test = train_test_split(
39     X_scaled, y, test_size=0.3, random_state=42
40 )
41
42 # Step 7: Unsupervised Learning - KMeans
43 kmeans = KMeans(n_clusters=3, random_state=42, n_init=10) # Assuming 3 classes
44 clusters = kmeans.fit_predict(X_scaled)
45
46 # Visualize clusters (using u_g and g_r)
47 plt.figure(figsize=(8, 6))
48 plt.scatter(X['u_g'], X['g_r'], c=clusters, cmap='viridis', s=50)
49 plt.title('KMeans Clustering (Unsupervised)')
50 plt.xlabel('u_g')
51 plt.ylabel('g_r')
52 plt.colorbar(label='Cluster')
53 plt.show()
54
55 # Step 8: Hyperparameter tuning for SVM
56 param_grid = {
57     'C': [0.1, 1, 10, 100],
58     'gamma': ['scale', 'auto', 0.001, 0.01, 0.1, 1]
59 }
60 grid = GridSearchCV(SVC(kernel='rbf', random_state=42), param_grid, cv=5, n_jobs=-1)
61 grid.fit(X_train, y_train)
62
63 print("Best parameters:", grid.best_params_)
64 print("Best cross-validation score:", grid.best_score_)
65
66 # Step 9: Train SVM with best parameters
```

```

67 svm_rbf = SVC(kernel='rbf', **grid.best_params_, random_state=42)
68 svm_rbf.fit(X_train, y_train)
69
70 # Step 10: Predict
71 y_train_pred = svm_rbf.predict(X_train)
72 y_test_pred = svm_rbf.predict(X_test)
73
74 # Step 11: Calculate accuracy
75 train_acc = accuracy_score(y_train, y_train_pred)
76 test_acc = accuracy_score(y_test, y_test_pred)
77 print(f"Training Accuracy (SVM-RBF): {train_acc:.2f}")
78 print(f"Test Accuracy (SVM-RBF): {test_acc:.2f}")
79
80 # Step 12: Predict class of a sample from test set
81 sample = X_test[[0]] # Use index-based selection for NumPy array
82 predicted_class = svm_rbf.predict(sample)[0]
83 print(f"Predicted Class for Sample: {predicted_class}")
84
85 # Step 13: Model fit evaluation
86 if train_acc > test_acc + 0.1:
87     model_fit = "Overfit"
88 elif abs(train_acc - test_acc) < 0.05 and train_acc > 0.9:
89     model_fit = "Best Fit"
90 else:
91     model_fit = "Underfit"
92 print(f"Model Fit: {model_fit}")

```



```

1 from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
2 import matplotlib.pyplot as plt
3
4 # 1. Compute the confusion matrix
5 cm = confusion_matrix(y_test, y_test_pred, labels=svm_rbf.classes_)
6
7 # 2. Display the confusion matrix
8 disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=svm_rbf.classes_)
9 plt.figure(figsize=(8, 6))
10 disp.plot(cmap='Blues', values_format='d')
11 plt.title("Confusion Matrix - SVM with RBF Kernel")
12 plt.show()
13

```